

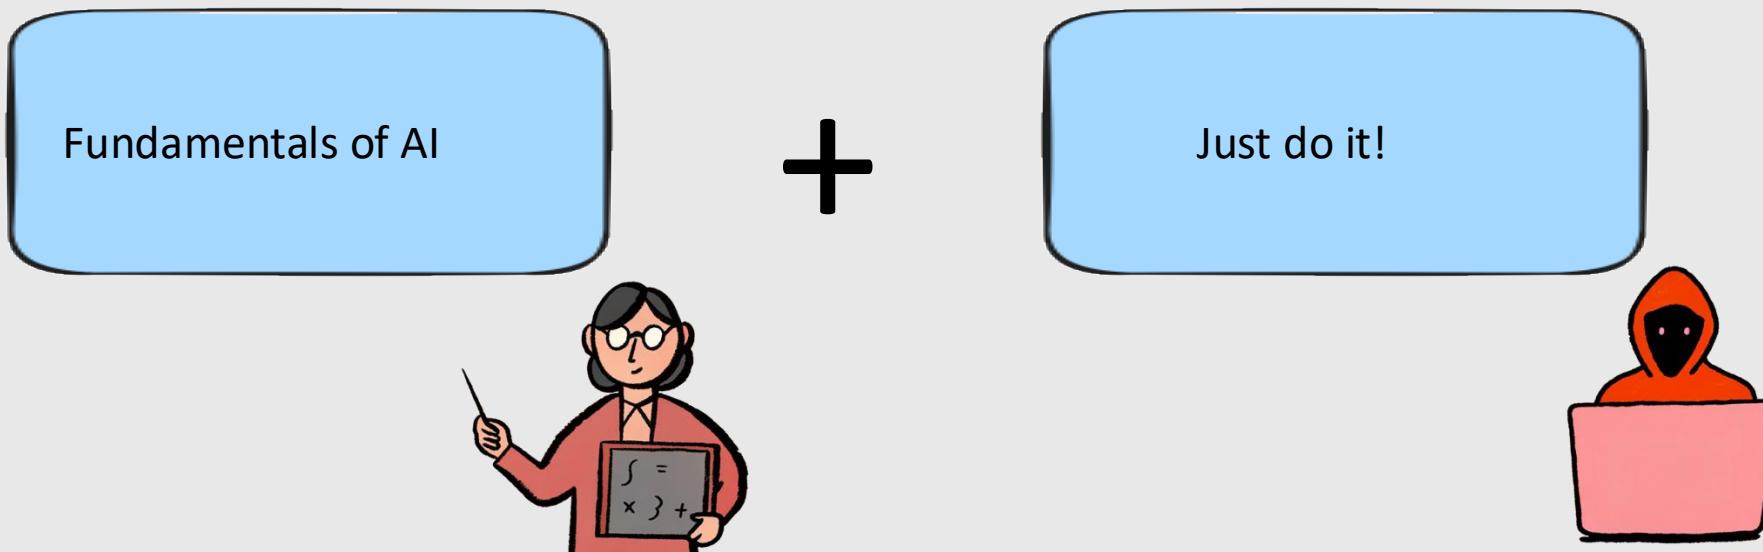
(tiny and tidy)

# Language Models

Leonardo Cotta, 2025

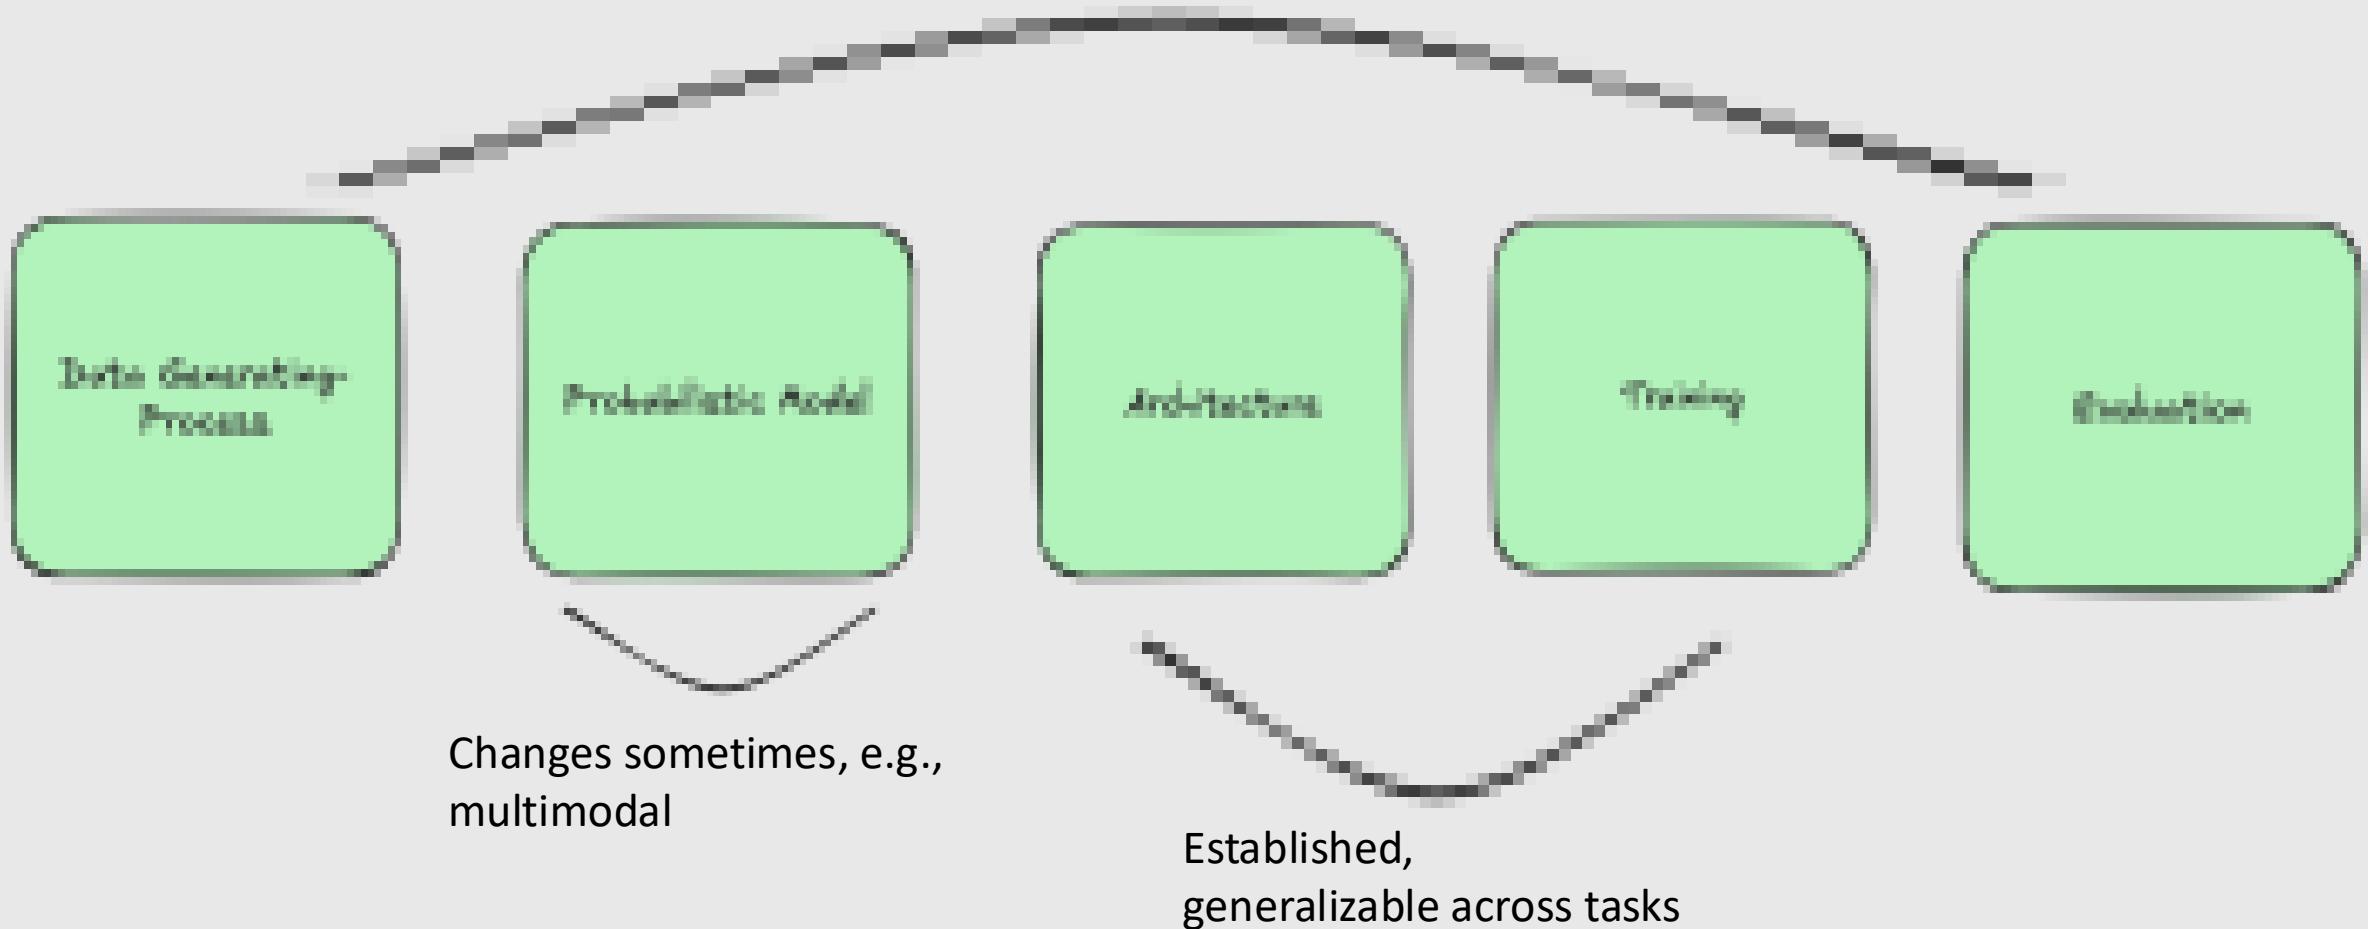
# A modern problem

- For most of us, it feels overwhelming to keep up with LLM progress...  
How can we do it?
- Our take in this two-day course:



# How to think about modern AI projects

Where we spend most of our time!



# Our focus: Pretraining

Due to time restrictions, we will focus on the pretraining part of modern language models and discuss how to learn and think about post-training, RL, etc. in the end!

## PART I:

- Data-generating process;
- Probabilistic Model.

# Language Data

- What's a language sample  $x$ ?
  - Something that somebody wrote down, i.e., "a document".
- What's the data-generating process of these docs?
  - People writing on Wikipedia, Shakespeare writing plays, or everyone writing on the Internet.
  - Each data-generating process induces a probability density of docs  $p(x)$ .
- How can we model  $p(x)$ ?

# Language Model

- Before we make our modeling choices, we need to specify more things about the random variable X.
- What's the support of X?
  - Sequence of characters? [l][ ][a][m][ ][a][ ][r][o][b][o][t].
  - $\text{Supp}(x) := V^*$  (product space).
- Important: What do we really want to do?
  - Sample from some approximation  $x \sim P_\theta(x)$
  - Sample from an arbitrary conditional  $y \sim P_\theta(y|x)$ 
    - Note: y is a completion of x.

# Language Model

- The magic of Bayes Theorem:
  - Auto-regressive models solve both problems!

$$p(x) = p(x_0) \prod_i p(x_i | x_0, \dots, x_{i-1})$$

[ $p(x_i | x_{-i})$ ]

- Mhm... what about  $p(x_0)$ ?
  - Let's just add a special symbol  $\langle \text{bos} \rangle$  to the vocab, and always start with it. Then, we don't need to learn the marginal, since we know  $p(x) = 1$ , if  $x = \langle \text{bos} \rangle$ , 0 otherwise.
  - Our goal is now to learn  $p(x)$  by approximating its conditionals  $p(x_i | x_{-i})$ !

# Learning

- A model for next-token (letter) prediction:

$$P_{\theta}(x) : V^* \rightarrow [O, I]^{|V|}$$

- How do we define the loss of the model over a single document  $x$ ?
  - Let's do it together, think of Maximum-Likelihood Estimation.

$$L(x, \theta) = -\log P_{\theta}(x)$$

- Now, with i.i.d. samples, we just average their losses.

# Learning

- We are ready to apply SGD and learn the model, as you learned in a previous NN course ;)
- Before we discuss what's the actual functional form of the model, let's discuss the following:
- Does our vocab need to contain characters?
  - Why are words and characters not used?
  - What tokenizers do you know?

## PART II:

- Tokenization.

# Why do we need tokenizers?

- Why do we think words are better than characters?
  - Or is it the other way around?
  - Both?
  - Why not bytes?
- Characters:
  - Documents become huge sequences;
  - Lot of redundancy.
- The distribution of words is usually quite skewed;
  - Some words are very rare, hard to learn on their own;
  - They can share information with others via sub-words!

# Tokenization and Compression

- Words decrease sequence size but reserve symbols in the vocab that we barely use.
  - We can have a unique symbol per example in the corpus, and not learn anything.
  - Larger vocab -> larger sample complexity.
- What do we want?
  - Balance sequence and vocab size.
- Idea:
  - Given a fixed vocab size, that we define;
  - Find the vocab that minimizes the average sequence length and appear often!

# BPE: The current winner

- Most modern LLMs use a very simple tokenization algorithm, which is also a compression algorithm from 1994: Byte-pair encoding!

For an N-size vocab:

1. Start with characters;
2. Repeatedly merge most frequent adjacent pair;
3. Stop after N merges.

## PART III:

- Transformers go brrr.

# Functional form

- Recall that we are now left with the task of designing a parametric function (prob density)

$$P_\theta(x) : V^* \rightarrow [0, 1]^{|V|}$$

- Input: context (prefix).
- Output: next token.

# Attention in a nutshell

- Attention is a function mapping a sequence of embeddings to another.

---

Algorithm I: Token embedding.

---

Input:  $v \in V \cong [N_V]$ , a token ID.

Output:  $e \in \mathbb{R}^{d_e}$ , the vector representation of the token.

Parameters:  $W_e \in \mathbb{R}^{d_e \times N_V}$ , the token embedding matrix

return  $e = W_e[:, v]$

---

# Attention in a nutshell

---

Algorithm I: Basic single-query attention.

---

Input:  $e \in \mathbb{R}^{d_{in}}$ , vector representation of the current token

Input:  $e_t \in \mathbb{R}^{d_{in}}$ , vector representations of context tokens  
 $t \in [T]$ .

Output:  $\tilde{v} \in \mathbb{R}^{d_{out}}$ , vector representation of the token and  
context combined.

Parameters:  $W_Q, W_k \in \mathbb{R}^{d_{attn} \times d_{in}}$ ,  $B_Q, B_k \in \mathbb{R}^{d_{attn}}$ , the query and  
key linear projections.

Parameters:  $W_v \in \mathbb{R}^{d_{out} \times d_{in}}$ ,  $B_v \in \mathbb{R}^{d_{out}}$ , the value linear  
projection.

$$Q \leftarrow W_Q e + B_Q;$$

$$\forall t : k_t \leftarrow W_k e_t + B_k;$$

$$\forall t : v_t \leftarrow W_v e_t + B_v;$$

$$\forall t : \alpha_t = \frac{\exp(Q^T k_t / \sqrt{d_{attn}})}{\sum_u \exp(Q^T k_u / \sqrt{d_{attn}})};$$

$$\text{return } \tilde{v} = \sum_{t=1}^T \alpha_t v_t$$

---

# Attention in a nutshell

---

Algorithm 1:  $\hat{V} \leftarrow \text{Attention}(X, Z | W_{QKV}, \text{Mask})$

---

// Computes a single (masked) self- or cross-attention head.

Input:  $X \in \mathbb{R}^{d_x \times l_x}$ ,  $Z \in \mathbb{R}^{d_z \times l_z}$ , vector representations of primary and context sequence.

Output:  $\hat{V} \in \mathbb{R}^{d_{\text{out}} \times l_x}$ , updated representations of tokens in  $X$ , folding in information from tokens in  $Z$ .

Parameters :  $W_{QKV}$  consisting of:

$$W_Q \in \mathbb{R}^{d_{\text{attn}} \times d_x}, B_Q \in \mathbb{R}^{d_{\text{attn}}}$$

$$W_K \in \mathbb{R}^{d_{\text{attn}} \times d_z}, B_K \in \mathbb{R}^{d_{\text{attn}}}$$

$$W_V \in \mathbb{R}^{d_{\text{out}} \times d_z}, B_V \in \mathbb{R}^{d_{\text{out}}}$$

Hyperparameters: Mask  $\in \{0, 1\}^{l_z \times l_x}, \uparrow(3)$

$$Q \leftarrow W_Q X + B_Q I^\top \quad [[\text{Query} \in \mathbb{R}^{d_{\text{attn}} \times l_x}]];$$

$$K \leftarrow W_K Z + B_K I^\top \quad [[\text{Key} \in \mathbb{R}^{d_{\text{attn}} \times l_z}]];$$

$$V \leftarrow W_V Z + B_V I^\top \quad [[\text{Value} \in \mathbb{R}^{d_{\text{out}} \times l_z}]];$$

$$S \leftarrow K^\top Q \quad [[\text{Score} \in \mathbb{R}^{l_z \times l_x}]];$$

$\forall t_z, t_x, \text{if } \neg \text{Mask}[t_z, t_x] \text{ then } S[t_z, t_x] \leftarrow -\infty;$

return  $\hat{V} = V \cdot \text{softmax}(S / \sqrt{d_{\text{attn}}})$

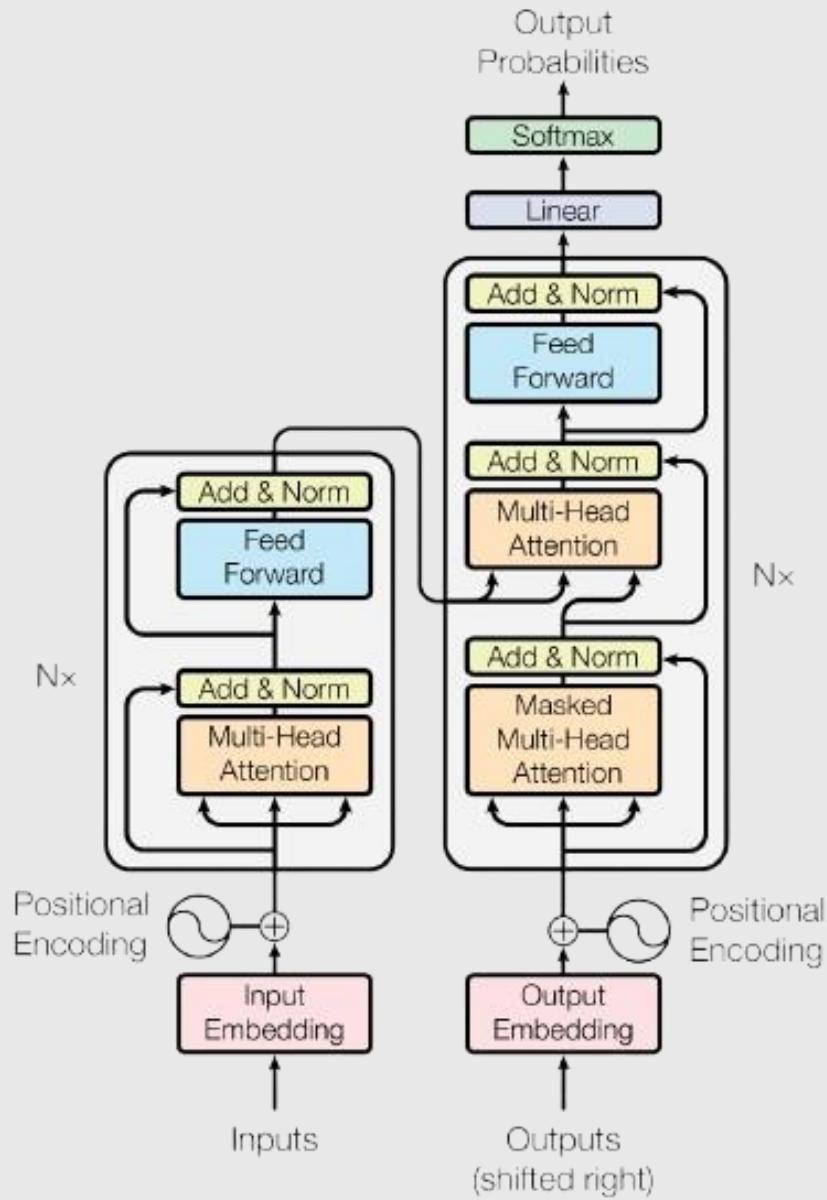
---

# Attention is almost all you need

- Attention is a linear mechanism that propagates information across a set of embeddings.
  - What's missing?
- Non-linearity?
- Positional information?
- How do we go to next-token distribution?
- We need the transformer!
  - Attention  $\neq$  transformer

# All you need (or not)

- How do we do these operations for a batch of sequences?
  - Let's do it together on the board!



# Positional Encodings

- In theory we could just have independent position embeddings, i.e. embedding 1, 2, 3, etc and combine them with the token embeddings.
  - Why is that a bad idea?
  - Sharing information is a deep learning secret.

# Sinusoidal encodings

For each position  $\text{pos} \in \{1, 2, \dots, n\}$  and dimension  $i \in \{0, 1, \dots, d-1\}$ , the positional encoding is defined as:

$$\text{PE}_{(\text{pos}, 2i)} = \sin \left( \frac{\text{pos}}{10000^{2i/d}} \right)$$

$$\text{PE}_{(\text{pos}, 2i+1)} = \cos \left( \frac{\text{pos}}{10000^{2i/d}} \right)$$

where:

- Even dimensions ( $2i$ ) use sine function
- Odd dimensions ( $2i+1$ ) use cosine function
- The wavelength forms a geometric progression from  $2\pi$  to  $10000 \cdot 2\pi$

# RoPE

- Rotation Matrix:

For position  $m$  and dimension pair  $(2i, 2i+1)$ , define the rotation angle:

$$\vartheta_i = 10000^{-2i/d}$$

The 2D rotation matrix for dimension pair  $i$  at position  $m$  is:

$$R_{\vartheta_i, m} = \begin{bmatrix} \cos(m\vartheta_i) & -\sin(m\vartheta_i) \\ \sin(m\vartheta_i) & \cos(m\vartheta_i) \end{bmatrix}$$

# RoPE

- Transformation:

For a query or key vector  $x_m \in \mathbb{R}^d$  at position  $m$ , partition it into pairs:

$$x_m = [x_m^{(1)}, x_m^{(2)}, \dots, x_m^{(d/2)}]^T$$

where each  $x_m^{(i)} = [x_{m,2i}, x_{m,2i+1}]^T \in \mathbb{R}^2$ .

Apply rotation to each pair:

$$\tilde{x}_m^{(i)} = R_{\vartheta_i, m} \cdot x_m^{(i)}$$

The full RoPE transformation is:

$$\text{ROPE}(x_m, m) = [\tilde{x}_m^{(1)}, \tilde{x}_m^{(2)}, \dots, \tilde{x}_m^{(d/2)}]^T$$

# RoPE

For each dimension pair  $(2i, 2i+1)$ :

$$\tilde{x}_{m,2i} = x_{m,2i} \cos(m\vartheta_i) - x_{m,2i+1} \sin(m\vartheta_i)$$

$$\tilde{x}_{m,2i+1} = x_{m,2i} \sin(m\vartheta_i) + x_{m,2i+1} \cos(m\vartheta_i)$$

# RoPE

Apply RoPE to queries and keys before computing attention:

$$Q = \text{RoPE}(xw_Q)$$

$$K = \text{RoPE}(xw_K)$$

$$V = xw_V$$

Then compute attention as usual:

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

## PART IV:

- Inference engine;
- Chat about advanced topics.

# How do we sample next tokens?

Let's look at engine.py



# RLHF is just Bayesian Inference

Let's do it together

